

# “Network Anomaly Detection for Intrusion Detection Systems Using Q-Learning and Deep Q-Learning: A Reinforcement Learning Approach”

2025 IEEE 9th International Conference on Software Engineering & Computer Systems (ICSECS)

Adhitama Wichaksono, Parman Sukaro, Aulia Arif Wardana

# Introduction

## Research Problem:

- Modern cybersecurity systems face increasingly complex threats
- Signature-based IDS ineffective against zero-day attacks
- Anomaly-based IDS struggle with adaptation to network changes

## Research Objective:

- Develop adaptive intrusion detection using Reinforcement Learning
- Compare Q-Learning and Deep Q-Learning (DQN) approaches
- Evaluate performance on zero-day attack detection

# Research Background

## Challenges in Network Security:

- Critical business processes depend on network systems
- Advanced attacks (DDoS, malware, zero-day) are increasing
- Traditional IDS limitations:
  - Rule-based systems ineffective against new attacks
  - High false positive rates hampering response times
  - Inability to adapt to changing network patterns

## Reinforcement Learning Potential:

- Enables agents to learn optimal detection strategies
- Adapts in real-time to new situations
- Reward-based learning suitable for cybersecurity applications

# Related Works

Previous Studies:	Research Gaps:
<ul style="list-style-type: none"><li>Modirrousta et al. (2022): DRL with CNN on UNSW-NB15 &amp; BoT-IoT datasets - 99.17% accuracy</li><li>Alavizadeh et al. (2022): Deep Q-Learning on NSL-KDD dataset - 78% accuracy</li><li>Tellache et al. (2024): Multi-agent RL-DQN on CIC-IDS-2017 dataset - 99% accuracy</li></ul>	<ul style="list-style-type: none"><li>Zero-day attack detection capabilities</li><li>Handling imbalanced data</li><li>Practical deployment considerations</li><li>Real-world effectiveness evaluation</li></ul>

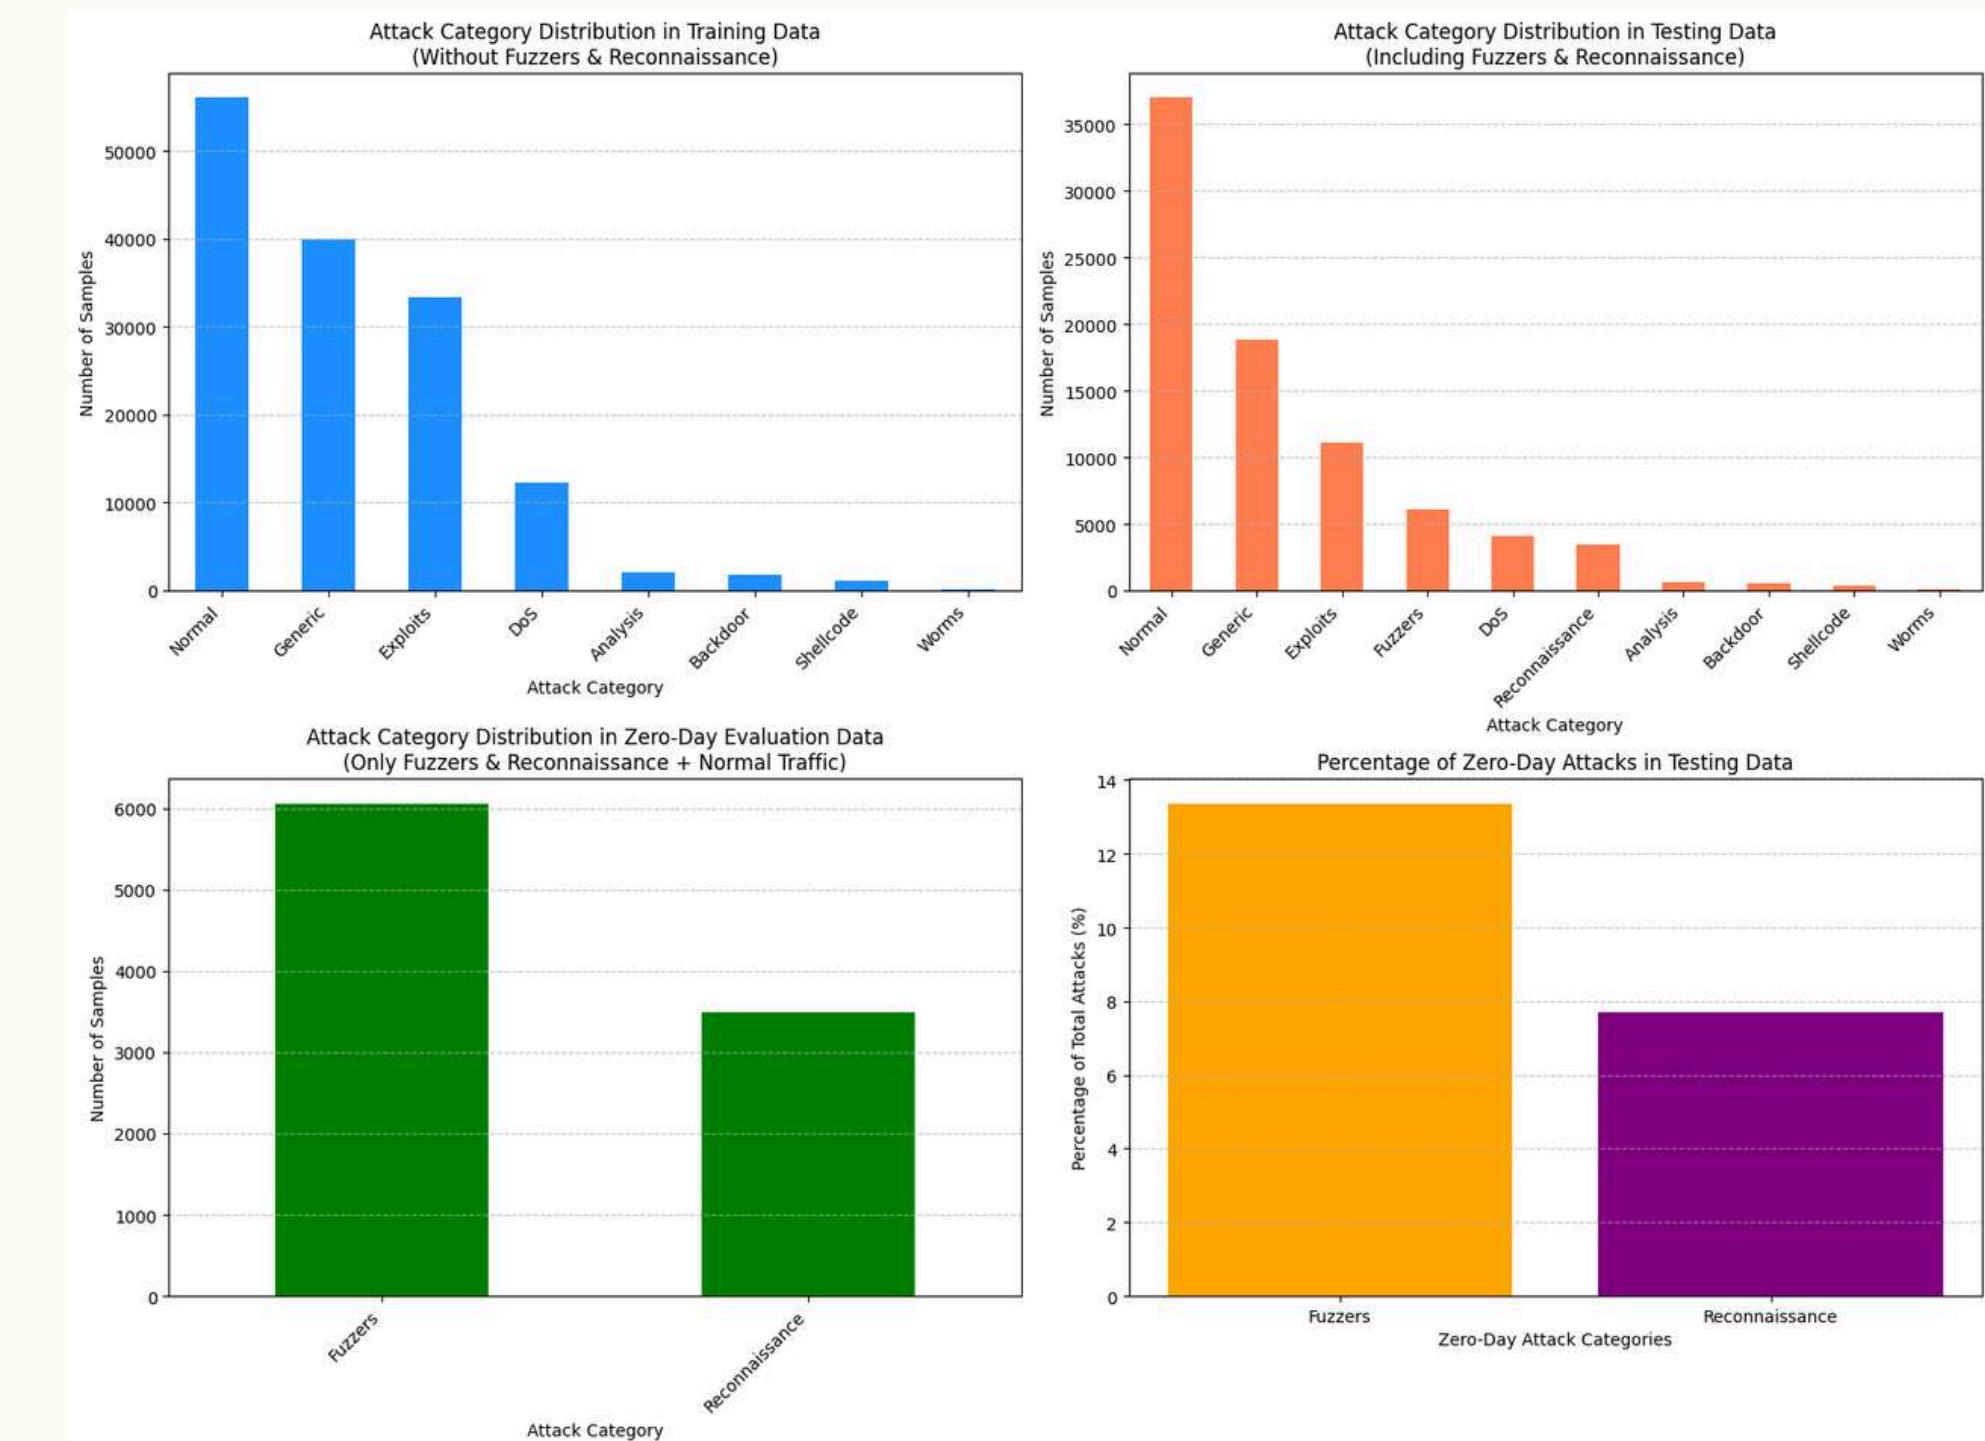
# Methodology Overview

Dataset:	Zero-Day Simulation Approach:
<ul style="list-style-type: none"><li>• UNSW-NB15 dataset from Australian Centre for Cyber Security</li><li>• 175,341 training records, 82,332 testing records</li><li>• 45 features (IP addresses, port numbers, protocols, traffic metrics)</li></ul>	<ul style="list-style-type: none"><li>• "Fuzzers" and "Reconnaissance" attack types removed from training</li><li>• Retained in test set to evaluate zero-day detection</li><li>• Created specialized zero-day evaluation dataset (46,558 samples)</li></ul>

# Data Preparation

## Dataset Distribution:

- Training data: Excludes "Fuzzers" & "Reconnaissance"
- Test data: Includes all attack types
- Zero-day evaluation: Contains only "Fuzzers", "Reconnaissance", and normal traffic



# Data Preprocessing

## Preprocessing Pipeline:

- Mean imputation for missing values
- Duplicate record removal
- Label encoding for categorical features
- SMOTE+Tomek for class imbalance (1:1 ratio)
- StandardScaler normalization
- Dimensionality reduction & feature selection

## Selected Features (20):

proto	sttl	djit	dmean
service	sload	swin	ct srv src
state	dload	stcpb	ct state ttl
dbytes	dinpkt	dtcpb	ct src ltm
rate	sjit	smean	is sm ips ports

# Model Development: Q-Learning

## Q-Learning Approach:

- Model-free reinforcement learning algorithm
- Uses Q-table to store expected utility of actions

Updates Q-values based on rewards:

$$Q(s, a) \leftarrow Q(s, a) + \alpha(r + \gamma \max_{a'} Q^*(s', a') - Q(s, a))$$

## Implementation Details:

- 30 states representing network traffic patterns
- Binary action space (2 actions): normal or anomalous
- Hyperparameter tuning with 5 different schemes

# Model Development: Deep Q-Learning

## DQN Approach:

- Neural network replaces Q-table for Q-function approximation
- Experience replay buffer for stable learning
- Target network updated periodically

## Network Architecture:

- Input layer (20 features)
- Two hidden layers with ReLU activation
- Output layer (2 actions)
- Trained to minimize loss:

$$L_i(\theta_i) = \mathbb{E}_{(s,a,r,s') \sim U(D)} \left[ \left( r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right)^2 \right]$$

# Experimental Setup

## Dataset Distribution:

### Training

- 86,709 samples (50%),
- 50.0% normal,
- 50.0% attacks (Analysis, Backdoor, DoS, Exploits, Generic, Shellcode, Worms)

### General Testing

- 82,332 samples
- 56.0% normal
- 44.0% attacks (all categories)

### Zero-day Testing

- 46,558 samples,
- 62.5% normal,
- 37.5% attacks (only Fuzzers & Reconnaissance + Normal)

# Hyperparameter Tuning:

To optimize the performance of the Q-Learning and DQN models, hyperparameter tuning was conducted through five different schemes.

Q-Learning Hyperparameter Tuning						DQN Hyperparameter Tuning				
Scheme	Learning Rate ( $\alpha$ )	Discount Factor ( $\gamma$ )	Initial Epsilon ( $\varepsilon$ )	Epsilon Decay	Minimum Epsilon	Scheme	Learning Rate	Epsilon Decay	Batch Size	Buffer Size
1	0.01	0.8	1.0	0.8	0.05	1	0.0001	0.93	32	50000
2	0.1	0.95	1.0	0.95	0.2	2	0.0003	0.95	64	50000
3	0.03	0.99	0.8	0.99	0.3	3	0.0005	0.97	128	50000
4	0.5	0.9	0.7	0.9	0.1	4	0.001	0.98	256	50000
5	0.2	0.7	0.9	0.85	0.01	5	0.003	0.99	512	50000

# Results: Q-Learning Performance on Standard Test Data

## Key Observations:

- Scheme 1 achieved the best overall performance with highest accuracy and F1 score
- Low learning rate (0.01) with moderate discount factor (0.8) performed best
- All schemes achieved high precision (>97%), showing minimal false positives
- Recall varied more significantly (86-90%), indicating challenges in detecting all attacks

Scheme	Accuracy	Precision	Recall	F1 Score	FPR	AUC
1	0.9418	0.9903	0.9030	0.9447	0.0108	0.9825
2	0.9259	0.9950	0.8699	0.9282	0.0054	0.9690
3	0.9235	0.9943	0.8661	0.9257	0.0061	0.9727
4	0.9181	0.9760	0.8727	0.9215	0.0262	0.9577
5	0.9235	0.9788	0.8801	0.9268	0.0233	0.9666

# Results: DQN Performance on Standard Test Data

## Key Observations:

- All DQN schemes significantly outperformed Q-Learning
- Scheme 5 achieved the highest accuracy (99.70%) and F1 score (0.9972)
- Scheme 1 had exceptional recall (99.98%), detecting virtually all attacks
- All schemes demonstrated exceptional AUC values (>0.989)
- FPR remained low across all configurations (0.36-1.99%)

Scheme	Accuracy	Precision	Recall	F1 Score	FPR	AUC
1	0.9909	0.9840	0.9998	0.9918	0.0199	0.9999
2	0.9888	0.9941	0.9855	0.9898	0.0071	0.9898
3	0.9958	0.9971	0.9954	0.9962	0.0036	0.9973
4	0.9945	0.9935	0.9965	0.9950	0.0079	0.9991
5	0.9970	0.9969	0.9975	0.9972	0.0038	0.9994

# Comparison of Best Q-Learning and DQN Schemes: Zero-Day Attack Detection Results

DQN misclassified only 9 attack instances as normal, while Q-Learning misclassified 3,387 attacks.

Performance on Unseen Attacks (Fuzzers & Reconnaissance):							Confusion Matrices for Zero-Day Attacks:				
Model	Accuracy	Precision	Recall	F1 Score	FPR	AUC	Model	TP	FP	TN	FN
Q-Learning	0.9187	0.9393	0.6456	0.7653	0.0108	0.9472	Q-Learning	6171	399	36601	3387
DQN	0.9840	0.9283	0.9991	0.9624	0.0199	0.9994	DQN	9549	737	36263	9

# Comparison with Traditional ML Models

Both reinforcement learning approaches outperformed traditional ML models, with DQN showing dramatically better zero-day detection than all alternatives.

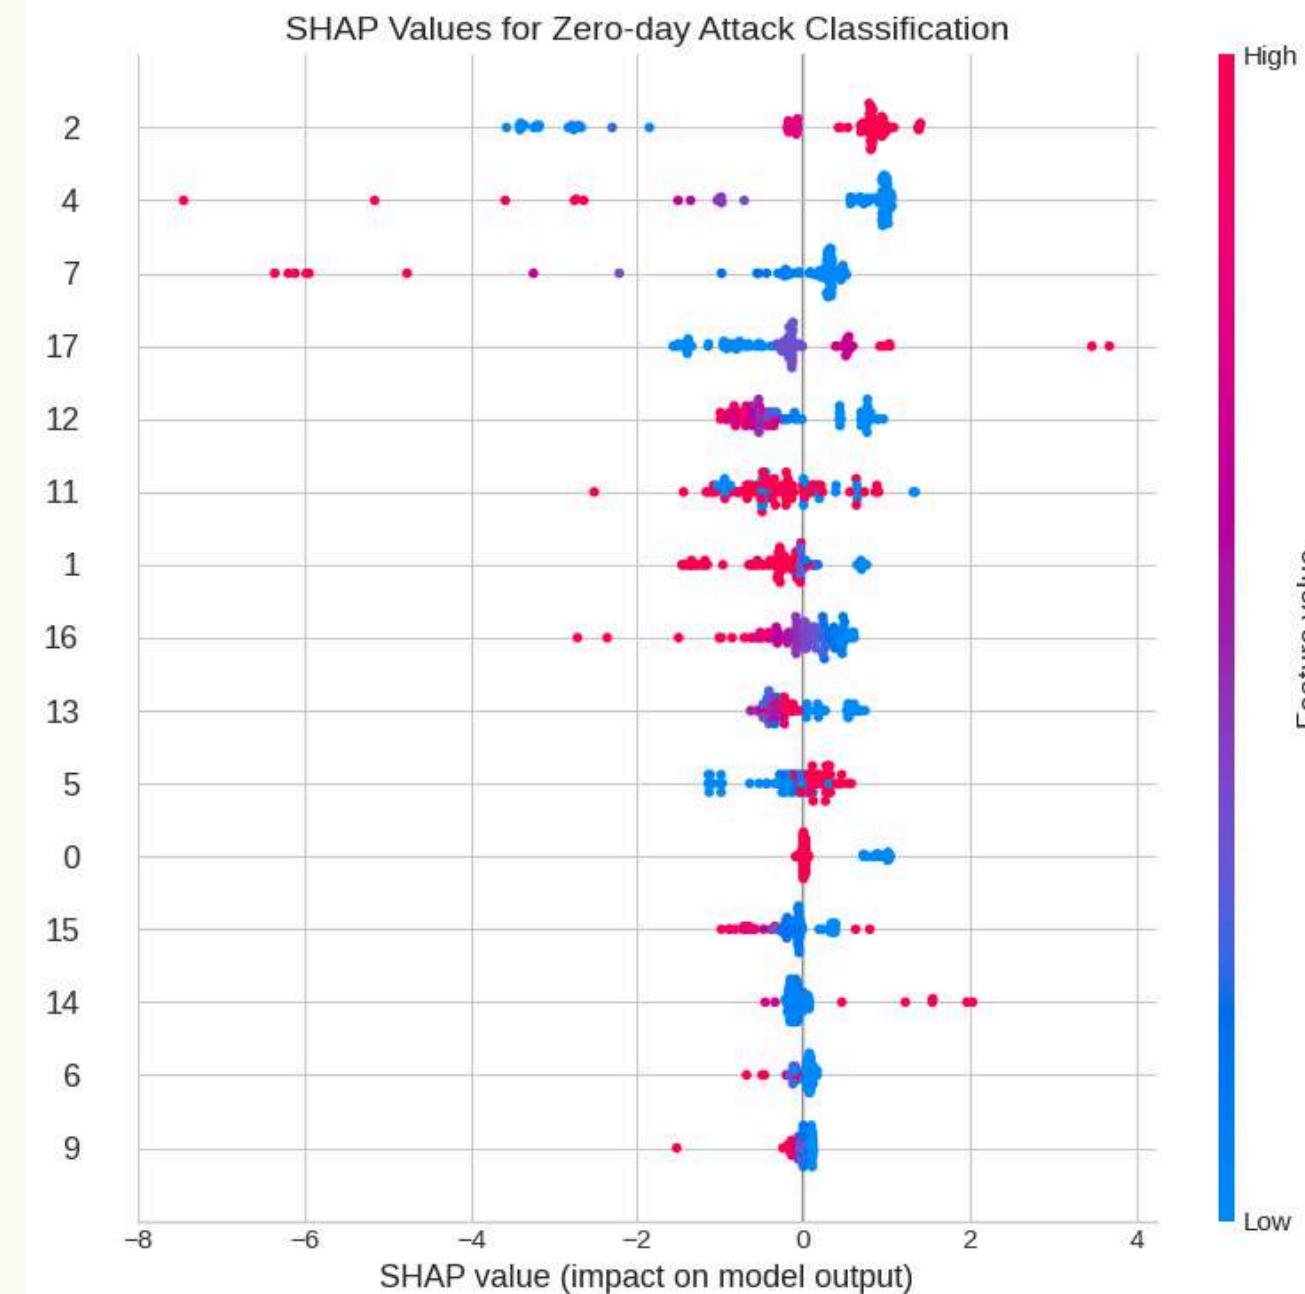
Performance on Standard Test Data:					Performance on Zero-Day Data:				
Metric	Q-Learning	DQN	Random Forest	SVM	Metric	Q-Learning	DQN	Random Forest	SVM
Accuracy	0.9418	0.9909	0.8978	0.8740	Accuracy	0.9187	0.9840	0.8360	0.8063
Precision	0.9903	0.9840	0.9462	0.9065	Precision	0.9393	0.9283	0.6508	0.5314
Recall	0.9030	0.9998	0.8634	0.8598	Recall	0.6456	0.9991	0.4337	0.4768
F1 Score	0.9447	0.9918	0.9029	0.8826	F1 Score	0.7653	0.9624	0.5205	0.5026
FPR	0.0108	0.0199	0.0601	0.1086	FPR	0.0108	0.0199	0.0601	0.1086
AUC	0.9825	0.9999	0.9667	0.9337	AUC	0.9472	0.9994	0.8605	0.7588

# SHAP Values for Zero-day Classification

## Key Insights:

- Service feature (2) shows strongest impact for distinguishing zero-day attacks
- Connection metrics (ct\_srv\_src) clearly separate normal from malicious traffic
- Data bytes (dbytes) contribute significantly to detection capability
- The visualization confirms effectiveness of feature selection
- Model demonstrates adaptability to previously unseen attack vectors
- Features with highest SHAP values align with cybersecurity domain knowledge

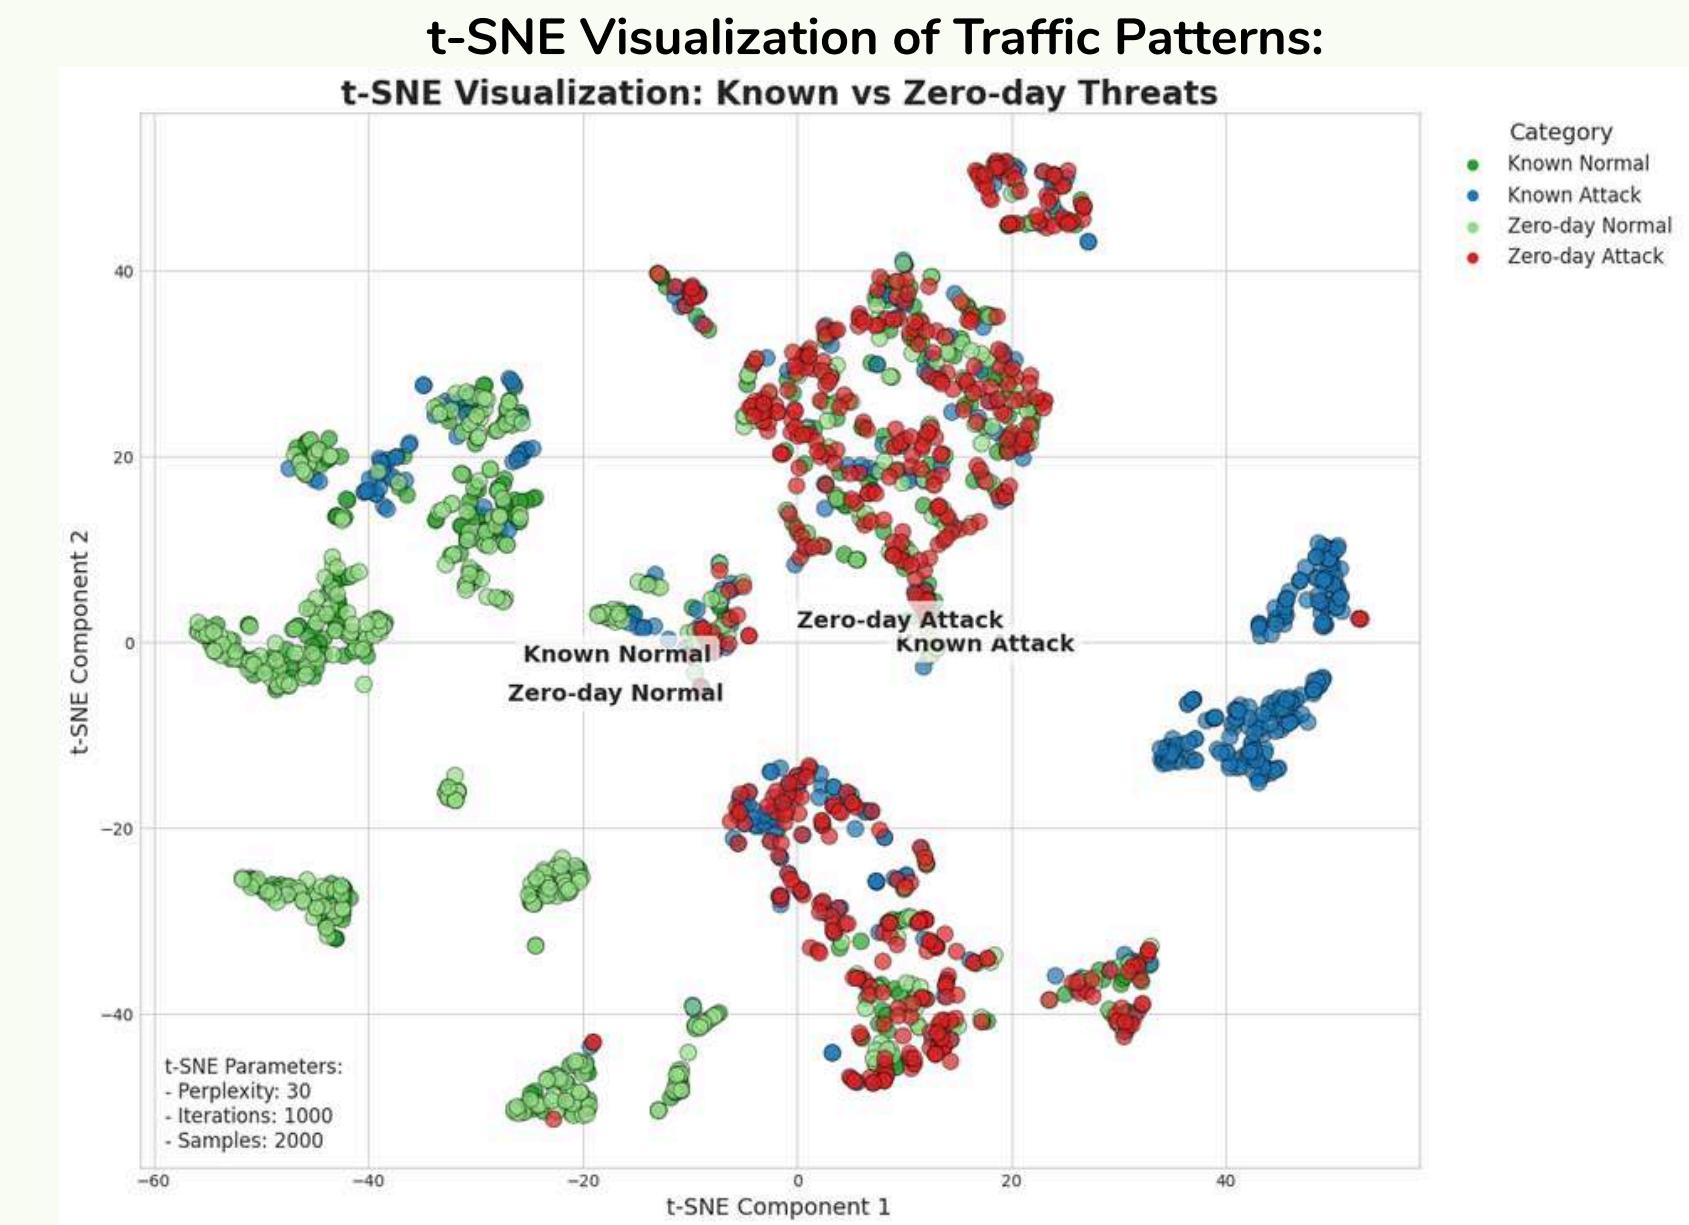
Detailed SHAP Analysis:



# t-SNE Visualization

## Visualization Insights:

- Four distinct clusters visible (Known Normal, Known Attack, Zero-day Normal, Zero-day Attack)
- DQN creates clear boundaries between normal and attack traffic
- Zero-day attack clusters positioned between known attacks and normal traffic
- Model learned generalized concept of "attackness" beyond specific signatures
- Known attacks form dense, well-defined clusters
- Zero-day attacks display more distributed patterns that partially overlap with known attack regions
- Effective transformation of feature space enhances attack-related patterns



# Computational Efficiency Comparison

Performance Metrics:					Efficiency Trade-offs:
Model	Training Time (sec)	Memory Usage (MB)	Model Size (MB)	Inference Speed (samples/sec)	
Q-Learning	36.03	1,247.39	0.000458	997,931	<ul style="list-style-type: none"> <li>• Q-Learning required only 0.74% of DQN's training time</li> <li>• Q-Learning used 48.5% of DQN's memory</li> <li>• Q-Learning model size dramatically smaller (0.000458 MB vs 0.03 MB)</li> <li>• Q-Learning inference speed approximately 275 times faster</li> <li>• These advantages must be weighed against DQN's superior detection capabilities</li> <li>• Computational demands may limit DQN deployment in resource-constrained environments</li> </ul>
DQN	4,852.5	2,572.1	0.03	3,623.4	

<b>CONCLUSION</b>	<b>FUTURE WORK</b>
<ul style="list-style-type: none"><li>• DQN proves to be more accurate in detecting network intrusions, achieving 99.09% accuracy compared to Q-Learning's 94.18%. DQN's advantage is most evident when facing new attacks (zero-day attacks) - its performance barely drops, while Q-Learning plummets to 64.56%.</li><li>• However, Q-Learning has its perks: it's lighter and faster. Training time is shorter, requires less memory, and the model size is smaller.</li></ul>	<p>Next Research Directions:</p> <ul style="list-style-type: none"><li>• Develop hybrid models that combine DQN's detection accuracy with Q-Learning's computational efficiency</li><li>• Test and optimize real-time detection performance in live production environments</li></ul>

# Thank you!

adhitamaw@student.telkomuniversity.ac.id

psukarno@telkomuniversity.ac.id

auliawardan@telkomuniversity.ac.id

Code Repository: <https://github.com/adhitamaw/Anomaly-Detection-using-Reinforcement-Learning>